


```
In [1]: #Import Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
import xgboost as xgb
```

```
In [2]: #Load Data
df = pd.read_csv('C:\liver_cirrhosis.csv')
df.head()
```

```
Out[2]:
```

	N_Days	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	Edema	Bilir
0	2221	C	Placebo	18499	F	N	Y	N	N	
1	1230	C	Placebo	19724	M	Y	N	Y	N	
2	4184	C	Placebo	11839	F	N	N	N	N	
3	2090	D	Placebo	16467	F	N	N	N	N	
4	2105	D	Placebo	21699	F	N	Y	N	N	

◀  ▶

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25000 entries, 0 to 24999
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   N_Days                25000 non-null  int64
1   Status                25000 non-null  object
2   Drug                 25000 non-null  object
3   Age                  25000 non-null  int64
4   Sex                  25000 non-null  object
5   Ascites              25000 non-null  object
6   Hepatomegaly         25000 non-null  object
7   Spiders              25000 non-null  object
8   Edema                25000 non-null  object
9   Bilirubin            25000 non-null  float64
10  Cholesterol           25000 non-null  float64
11  Albumin              25000 non-null  float64
12  Copper               25000 non-null  float64
13  Alk_Phos             25000 non-null  float64
14  SGOT                 25000 non-null  float64
15  Tryglicerides        25000 non-null  float64
16  Platelets            25000 non-null  float64
17  Prothrombin          25000 non-null  float64
18  Stage                25000 non-null  int64
dtypes: float64(9), int64(3), object(7)
memory usage: 3.6+ MB
```

```
In [4]: print(df.describe())
```

	N_Days	Age	Bilirubin	Cholesterol	Albumin \
count	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000
mean	1887.117040	18495.877080	3.402644	372.331471	3.486578
std	1091.690918	3737.596616	4.707491	193.668452	0.380488
min	41.000000	9598.000000	0.300000	120.000000	1.960000
25%	1080.000000	15694.000000	0.800000	275.000000	3.290000
50%	1680.000000	18499.000000	1.300000	369.510563	3.510000
75%	2576.000000	20955.000000	3.400000	369.510563	3.750000
max	4795.000000	28650.000000	28.000000	1775.000000	4.640000

	Copper	Alk_Phos	SGOT	Tryglicerides	Platelets \
count	25000.000000	25000.000000	25000.000000	25000.000000	25000.000000
mean	100.184663	1995.675597	123.166345	123.822548	256.007337
std	73.184840	1798.885660	47.747616	52.786350	98.679006
min	4.000000	289.000000	26.350000	33.000000	62.000000
25%	52.000000	1032.000000	92.000000	92.000000	189.000000
50%	97.648387	1828.000000	122.556346	124.702128	251.000000
75%	107.000000	1982.655769	134.850000	127.000000	311.000000
max	588.000000	13862.400000	457.250000	598.000000	721.000000

	Prothrombin	Stage
count	25000.000000	25000.000000
mean	10.734549	2.00116
std	0.904436	0.81387
min	9.000000	1.00000
25%	10.000000	1.00000
50%	10.600000	2.00000
75%	11.100000	3.00000
max	18.000000	3.00000

```
In [5]: print(df.isnull().sum())
```

```
N_Days      0
Status      0
Drug        0
Age         0
Sex         0
Ascites     0
Hepatomegaly 0
Spiders     0
Edema       0
Bilirubin   0
Cholesterol 0
Albumin     0
Copper      0
Alk_Phos    0
SGOT        0
Tryglicerides 0
Platelets   0
Prothrombin 0
Stage       0
dtype: int64
```

```
In [6]: print("\nClass Distribution:")
        print(df['Stage'].value_counts())
```


```
Class Distribution:
Stage
2      8441
3      8294
1      8265
Name: count, dtype: int64
```

```
In [7]: df = df.drop(columns=['N_Days'], errors='ignore')
        df
```

Out[7]:

	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	Edema	Bilirubin
0	C	Placebo	18499	F	N		Y	N	N
1	C	Placebo	19724	M	Y		N	Y	N
2	C	Placebo	11839	F	N		N	N	N
3	D	Placebo	16467	F	N		N	N	N
4	D	Placebo	21699	F	N		Y	N	N
...	...	...	...	...	...		...	...	...
24995	D	D-penicillamine	23612	F	N		N	N	N
24996	D	D-penicillamine	23612	F	N		N	N	N
24997	D	D-penicillamine	16736	F	N		Y	Y	Y
24998	C	D-penicillamine	16990	F	N		Y	N	N
24999	C	D-penicillamine	16990	F	N		Y	N	N

25000 rows × 18 columns

In [8]: `df['Age'].describe()`

```
Out[8]: count    25000.000000
mean      18495.877080
std        3737.596616
min         9598.000000
25%       15694.000000
50%       18499.000000
75%       20955.000000
max       28650.000000
Name: Age, dtype: float64
```

In [9]: `print(df['Age'].head())`

```
0    18499
1    19724
2    11839
3    16467
4    21699
Name: Age, dtype: int64
```

In [10]: `df['Age'] = df['Age']/365`

```
In [11]: #Encode Categorical Features
from sklearn.preprocessing import LabelEncoder

label_cols = ['Sex', 'Ascites', 'Hepatomegaly', 'Spiders', 'Edema', 'Drug', 'Status']
for col in label_cols:
    df[col] = LabelEncoder().fit_transform(df[col])
```

```
In [12]: print(df.head())
```

	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	Edema	\
0	0	1	50.682192	0	0	1	0	0	
1	0	1	54.038356	1	1	0	1	0	
2	0	1	32.435616	0	0	0	0	0	
3	2	1	45.115068	0	0	0	0	0	
4	2	1	59.449315	0	0	1	0	0	

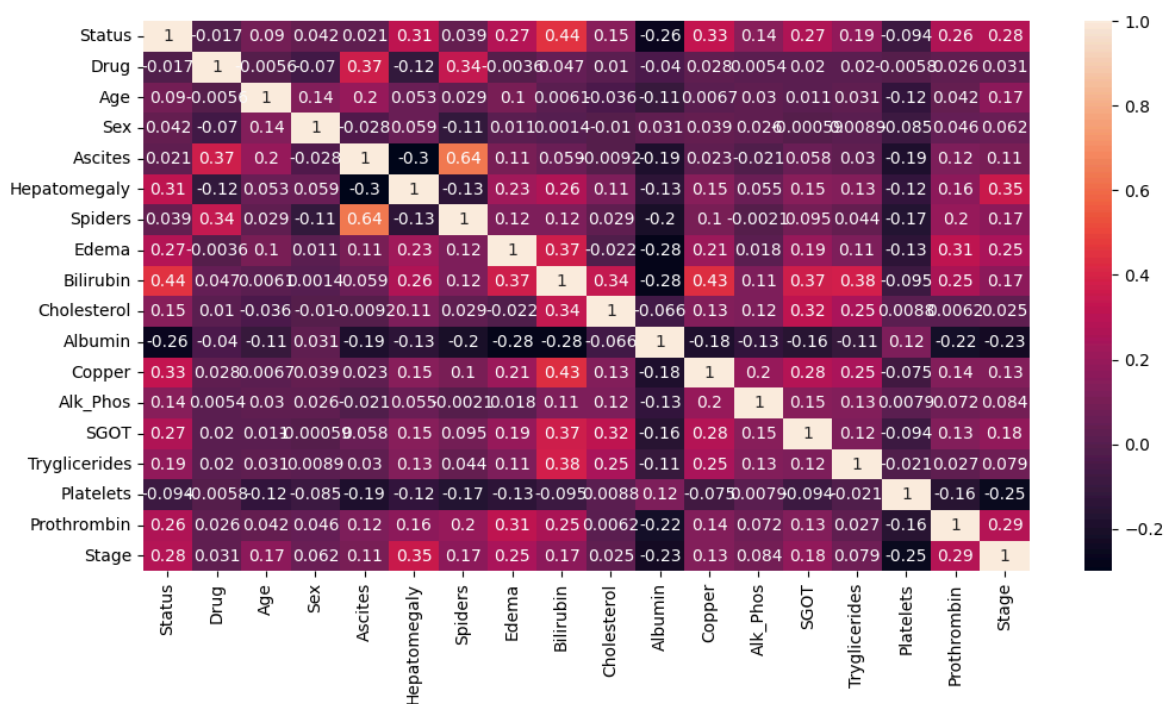
	Bilirubin	Cholesterol	Albumin	Copper	Alk_Phos	SGOT	Tryglicerides	\
0	0.5	149.0	4.04	227.0	598.0	52.70	57.0	
1	0.5	219.0	3.93	22.0	663.0	45.00	75.0	
2	0.5	320.0	3.54	51.0	1243.0	122.45	80.0	
3	0.7	255.0	3.74	23.0	1024.0	77.50	58.0	
4	1.9	486.0	3.54	74.0	1052.0	108.50	109.0	

	Platelets	Prothrombin	Stage
0	256.0	9.9	1
1	220.0	10.8	2
2	225.0	10.0	2
3	151.0	10.2	2
4	151.0	11.5	1

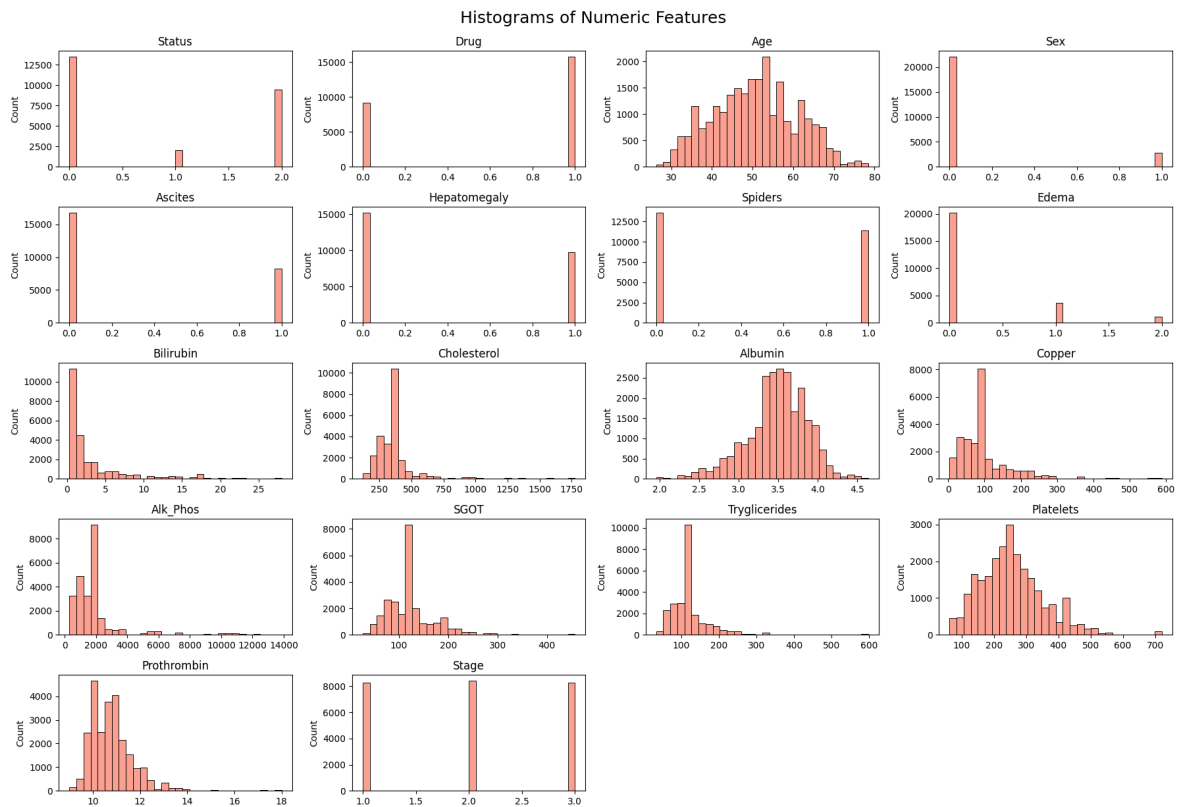
```
In [13]: plt.figure(figsize=(12,6))
sns.heatmap(df.corr(), annot=True)
```

Out[13]: <Axes: >



```
In [14]: numeric_columns = df.select_dtypes(include='number').columns
plt.figure(figsize=(18, 12))
for i, col in enumerate(numeric_columns, 1):
```

```
plt.subplot((len(numeric_columns) - 1) // 4 + 1, 4, i)
sns.histplot(df[col], bins=30, color='salmon', edgecolor='black')
plt.title(col)
plt.xlabel('')
plt.ylabel('Count')
plt.tight_layout()
plt.suptitle("Histograms of Numeric Features", fontsize=18, y=1.02)
plt.show()
```



```
In [15]: #Recode Target Variable (Stage)
df['Stage'] = df['Stage'].replace({1:0,2:1,3:2})
```

```
In [16]: df
```

Out[16]:

	Status	Drug	Age	Sex	Ascites	Hepatomegaly	Spiders	Edema	Bilirubin
0	0	1	50.682192	0	0	1	0	0	0.5
1	0	1	54.038356	1	1	0	1	0	0.5
2	0	1	32.435616	0	0	0	0	0	0.5
3	2	1	45.115068	0	0	0	0	0	0.7
4	2	1	59.449315	0	0	1	0	0	1.5
...	...	...	...	...	...	...	...	...	...
24995	2	0	64.690411	0	0	0	0	0	0.8
24996	2	0	64.690411	0	0	0	0	0	0.8
24997	2	0	45.852055	0	0	1	1	2	5.1
24998	0	0	46.547945	0	0	1	0	0	0.8
24999	0	0	46.547945	0	0	1	0	0	0.8

25000 rows × 18 columns

In [17]: `df['Stage']`

Out[17]:

```

0      0
1      1
2      1
3      1
4      0
..
24995  1
24996  1
24997  2
24998  1
24999  1
Name: Stage, Length: 25000, dtype: int64

```

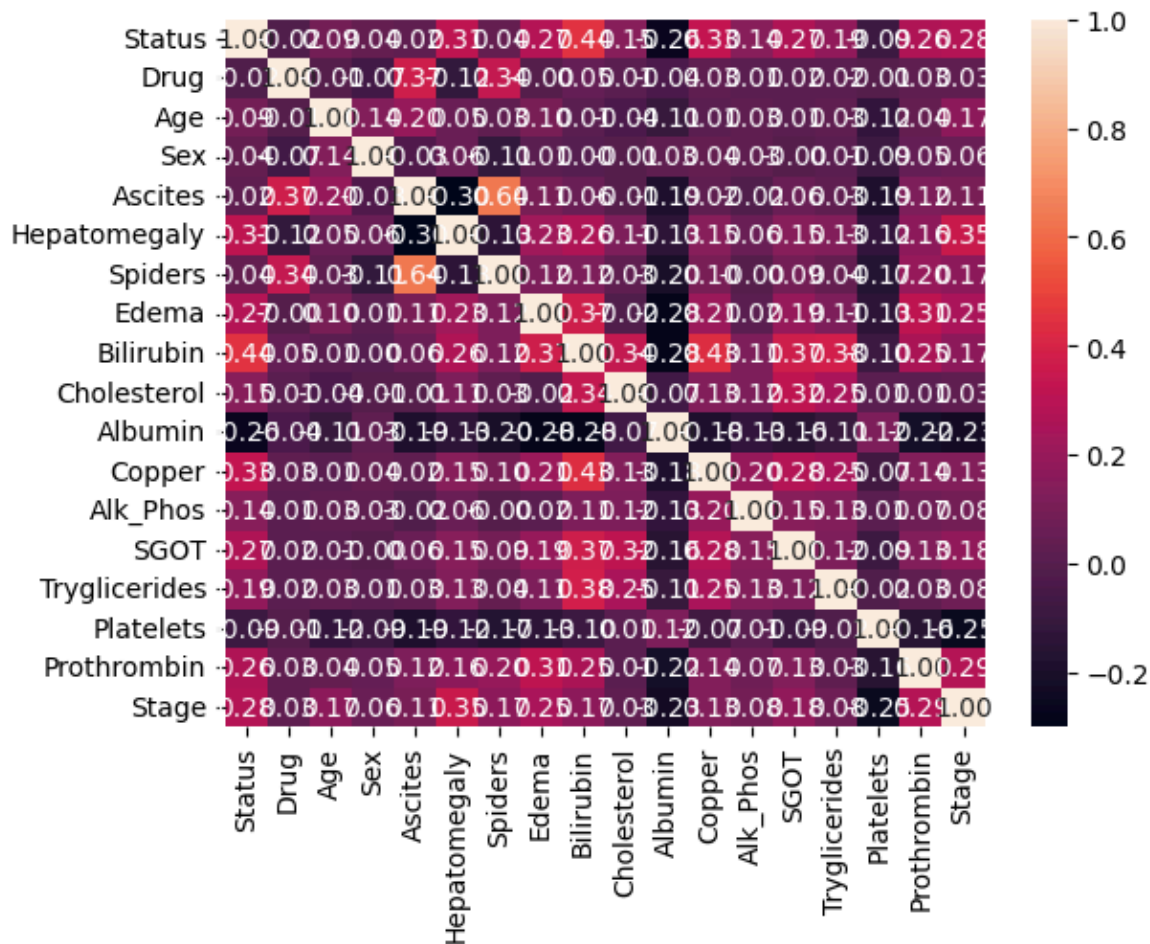
In [18]:

```

#Exploratory Data Analysis
sns.countplot(x='Status', data=df)
sns.heatmap(df.corr(), annot=True, fmt='.2f')

```

Out[18]: &lt;Axes: &gt;



```
In [19]: # Feature and Target Split
X = df.drop(columns=['Status'])
y = df['Status']
```

```
In [20]: #Scale Features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [21]: #Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, r
```

```
In [22]: #Train Multiple Models
log_reg = LogisticRegression(max_iter=1000)
log_reg.fit(X_train, y_train)
y_pred_log_reg = log_reg.predict(X_test)

print("Logistic Regression Performance")
print(confusion_matrix(y_test, y_pred_log_reg))
print(classification_report(y_test, y_pred_log_reg))
```



## Logistic Regression Performance

```
[[2381    1   321]
 [ 212   18   176]
 [ 770    7 1114]]
```

	precision	recall	f1-score	support
0	0.71	0.88	0.79	2703
1	0.69	0.04	0.08	406
2	0.69	0.59	0.64	1891
accuracy			0.70	5000
macro avg	0.70	0.50	0.50	5000
weighted avg	0.70	0.70	0.67	5000

```
In [23]: rf_clf = RandomForestClassifier(class_weight='balanced', random_state=42)
rf_clf.fit(X_train, y_train)
y_pred_rf = rf_clf.predict(X_test)

print("Random Forest Performance")
print(confusion_matrix(y_test, y_pred_rf))
print(classification_report(y_test, y_pred_rf))
```

## Random Forest Performance

```
[[2690    2    11]
 [  14  383    9]
 [  28    0 1863]]
```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	2703
1	0.99	0.94	0.97	406
2	0.99	0.99	0.99	1891
accuracy			0.99	5000
macro avg	0.99	0.97	0.98	5000
weighted avg	0.99	0.99	0.99	5000

```
In [24]: xgb_clf = xgb.XGBClassifier(eval_metric='mlogloss', use_label_encoder=False)
xgb_clf.fit(X_train, y_train)
y_pred_xgb = xgb_clf.predict(X_test)

print("XGBoost Performance")
print(confusion_matrix(y_test, y_pred_xgb))
print(classification_report(y_test, y_pred_xgb))
```

C:\Users\sakshi\AppData\Roaming\Python\Python311\site-packages\xgboost\training.py:  
183: UserWarning: [13:54:15] WARNING: C:\actions-runner\work\xgboost\xgboost\src\l  
earner.cc:738:

Parameters: { "use\_label\_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
```

## XGBoost Performance

```
[[2692   2    9]
 [   9 391    6]
 [  22   0 1869]]
```

	precision	recall	f1-score	support
0	0.99	1.00	0.99	2703
1	0.99	0.96	0.98	406
2	0.99	0.99	0.99	1891
accuracy			0.99	5000
macro avg	0.99	0.98	0.99	5000
weighted avg	0.99	0.99	0.99	5000

```
In [25]: svm_clf = SVC(class_weight='balanced', kernel='linear', random_state=42)
svm_clf.fit(X_train, y_train)
y_pred_svm = svm_clf.predict(X_test)

print("SVM Performance")
print(confusion_matrix(y_test, y_pred_svm))
print(classification_report(y_test, y_pred_svm))
```

## SVM Performance

```
[[2193  198  312]
 [   81  256   69]
 [  535  320 1036]]
```

	precision	recall	f1-score	support
0	0.78	0.81	0.80	2703
1	0.33	0.63	0.43	406
2	0.73	0.55	0.63	1891
accuracy			0.70	5000
macro avg	0.61	0.66	0.62	5000
weighted avg	0.73	0.70	0.70	5000

```
In [26]: #Evaluate Model Performance
models = ['Logistic Regression', 'Random Forest', 'XGBoost', 'SVM']
predictions = [y_pred_log_reg, y_pred_rf, y_pred_xgb, y_pred_svm]

for model, pred in zip(models, predictions):
    print(f"\n{model} Classification Report:")
    print(classification_report(y_test, pred))
    print(confusion_matrix(y_test, pred))
```

## Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.71	0.88	0.79	2703
1	0.69	0.04	0.08	406
2	0.69	0.59	0.64	1891
accuracy			0.70	5000
macro avg	0.70	0.50	0.50	5000
weighted avg	0.70	0.70	0.67	5000

```
[[2381  1 321]
 [ 212 18 176]
 [ 770  7 1114]]
```

## Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	2703
1	0.99	0.94	0.97	406
2	0.99	0.99	0.99	1891
accuracy			0.99	5000
macro avg	0.99	0.97	0.98	5000
weighted avg	0.99	0.99	0.99	5000

```
[[2690  2  11]
 [  14 383  9]
 [  28  0 1863]]
```

## XGBoost Classification Report:

	precision	recall	f1-score	support
0	0.99	1.00	0.99	2703
1	0.99	0.96	0.98	406
2	0.99	0.99	0.99	1891
accuracy			0.99	5000
macro avg	0.99	0.98	0.99	5000
weighted avg	0.99	0.99	0.99	5000

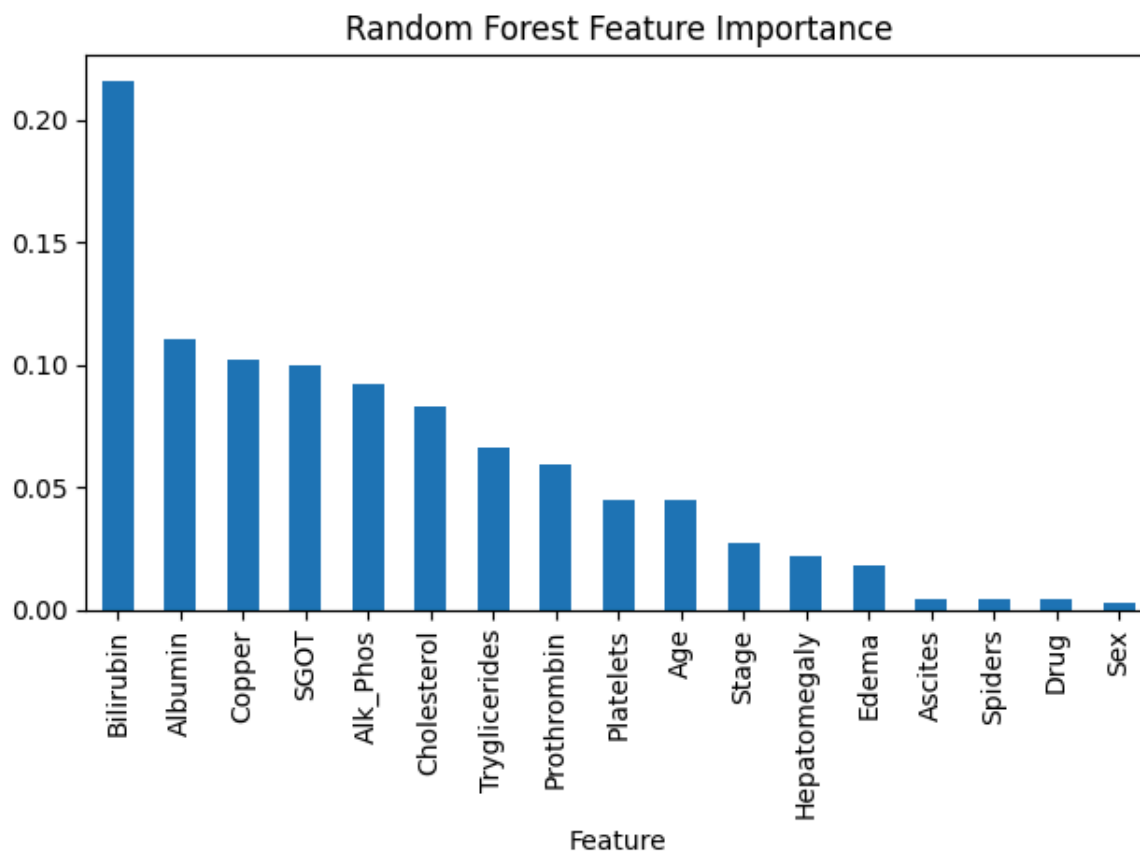
```
[[2692  2  9]
 [  9 391  6]
 [ 22  0 1869]]
```

## SVM Classification Report:


	precision	recall	f1-score	support
0	0.78	0.81	0.80	2703
1	0.33	0.63	0.43	406
2	0.73	0.55	0.63	1891
accuracy			0.70	5000
macro avg	0.61	0.66	0.62	5000
weighted avg	0.73	0.70	0.70	5000

```
[[2193 198 312]
 [  81 256  69]
 [ 535 320 1036]]
```

```
In [27]: #Feature Importance
importances_rf = rf_clf.feature_importances_
feat_df_rf = pd.DataFrame({'Feature': X.columns, 'Importance': importances_rf})
feat_df_rf = feat_df_rf.sort_values(by='Importance', ascending=False)
feat_df_rf.plot(kind='bar', x='Feature', y='Importance', legend=False)
plt.title('Random Forest Feature Importance')
plt.tight_layout()
plt.show()
```



```
In [28]: xgb.plot_importance(xgb_clf)
plt.title('XGBoost Feature Importance')
plt.tight_layout()
plt.show()
```



```
In [29]: #Save Best Model
import pickle
```

```
In [30]: with open('rf_model.pkl', 'wb') as f:
pickle.dump(rf_clf, f)
```

```
In [31]: with open('xgb_model.pkl', 'wb') as f:
pickle.dump(xgb_clf, f)
```

```
In [32]: #Cross-Validation
from sklearn.model_selection import cross_val_score
```

```
In [33]: log_reg_scores = cross_val_score(log_reg, X_scaled, y, cv=5, scoring='accuracy')
rf_scores = cross_val_score(rf_clf, X_scaled, y, cv=5, scoring='accuracy')
xgb_scores = cross_val_score(xgb_clf, X_scaled, y, cv=5, scoring='accuracy')
svm_scores = cross_val_score(svm_clf, X_scaled, y, cv=5, scoring='accuracy')

print("Logistic Regression CV Accuracy: ", log_reg_scores.mean())
print("Random Forest CV Accuracy: ", rf_scores.mean())
print("XGBoost CV Accuracy: ", xgb_scores.mean())
print("SVM CV Accuracy: ", svm_scores.mean())
```

```
C:\Users\sakshi\AppData\Roaming\Python\Python311\site-packages\xgboost\training.py:
183: UserWarning: [13:59:53] WARNING: C:\actions-runner\_work\xgboost\xgboost\src\l
earner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
C:\Users\sakshi\AppData\Roaming\Python\Python311\site-packages\xgboost\training.py:
183: UserWarning: [13:59:55] WARNING: C:\actions-runner\_work\xgboost\xgboost\src\l
earner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
C:\Users\sakshi\AppData\Roaming\Python\Python311\site-packages\xgboost\training.py:
183: UserWarning: [13:59:56] WARNING: C:\actions-runner\_work\xgboost\xgboost\src\l
earner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
C:\Users\sakshi\AppData\Roaming\Python\Python311\site-packages\xgboost\training.py:
183: UserWarning: [13:59:58] WARNING: C:\actions-runner\_work\xgboost\xgboost\src\l
earner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
C:\Users\sakshi\AppData\Roaming\Python\Python311\site-packages\xgboost\training.py:
183: UserWarning: [14:00:00] WARNING: C:\actions-runner\_work\xgboost\xgboost\src\l
earner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

Logistic Regression CV Accuracy: 0.69892

Random Forest CV Accuracy: 0.98612

XGBoost CV Accuracy: 0.99024


SVM CV Accuracy: 0.6759600000000001

```
In [34]: models = ['Logistic Regression', 'Random Forest', 'XGBoost', 'SVM']
cv_accuracies = [
    log_reg_scores.mean()*100,
    rf_scores.mean()*100,
    xgb_scores.mean()*100,
    svm_scores.mean()*100
]

plt.figure(figsize=(8, 5))
bars = plt.bar(models, cv_accuracies, color=['skyblue', 'lightgreen', 'orange', 'p

for bar, acc in zip(bars, cv_accuracies):
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 1,
             f'{acc:.2f}%', ha='center', va='bottom')

plt.ylim(0, 110)
plt.ylabel('Cross-Validated Accuracy (%)')
plt.title('Final Model Accuracy (5-Fold Cross-Validation)')
plt.grid(axis='y', linestyle='--', alpha=0.5)
plt.tight_layout()
plt.show()
```



In [ ]: